

# Artificial Intelligence II Homework 2

## Comments & Model Performance results on Question 3

Pavlos Spanoudakis (sdi1800184)

---

### Basic Execution flow

1. Reading train set & validation set data from the input files, into `DataFrame`'s. The file paths can be modified on the first notebook cell.
2. Checking if all the samples have the expected format without missing values.
3. Trimming the created `DataFrame`'s, keeping just the required data (tweets & labels).
4. Reading pre-trained word embeddings file, and storing the vector for each given word.
5. Vectorizing the train set & validation set samples using the pre-trained vectors:
  - We start by tokenizing each tweet using `nltk.TweetVectorizer`.
  - For each word/token returned by the tokenizer, we find the corresponding stored vector.
  - We sum up all the vectors and divide the result by the number of tokens in the tweet.  
This "mean vector" will be assigned to the tweet.  
If the vector for a word is not found, it does not contribute to the sum, but it does influence the final result, so unknown words essentially "contribute" with a vector of zeros.
6. Creating the Neural Network, selecting Loss function and Optimizer:
  - We create a `Network` instance, with the desired number and size of layers.
  - We create a `TensorDataset` in order create a `DataLoader` which will be used to obtain train set batches during each epoch.
  - We create X and Y Tensors for both sets, which will be used in multiple cases.
  - We select the Loss Function & the Optimizer.

Hyperparameters such as number of epochs, batch size, learning rate, number & size of layers etc. can be modified on cell #2.

7. Training the model:
8. Displaying performance results: After the end of training we display:
  - The Confusion Matrices of the final model predictions on the Train and Validation sets.
  - The F1 Learning Curves for both sets, to demonstrate the performance of the model after each epoch of the training phase.
  - The F1, Precision and Recall scores for the final model on both sets.  
We use the corresponding **scikit-learn** routines for calculating the scores. We use `average='micro'` for F1 score and `average='macro'` for Precision & Recall scores.
  - Training phase Loss curves for both sets.

- The ROC Curves for the Validation set predictions.  
We mirror the usage of `roc_curve` from **scikit-learn** for multiple classes, as demonstrated [here](#). To create the curves, we apply the `softmax` function to the NN output vector, to convert it to possibility values that add-up to 1, and use the `softmax` output to create the curves. `roc_curve` applies generated possibility thresholds to create the curves, therefore if we provided it with just the predicted labels, it would only apply 3 thresholds to each result, which is insufficient to create useful ROC curves.

## Different models performance comparison

For all the results below, `solver = "lbfgs"` was used in the Classifier.

- Using `TfidfVectorizer`:
  - Using `min_df`, `max_df` & `ngram_range` in the Vectorizer:

## Takeaways

- We see much better performance on **Neutral** and **Pro-Vaccine** tweets in all models, since a significant amount of train set tweets are labeled as such.  
Apparently, the majority of the **Anti-Vaccine** tweets are falsely predicted as **Pro-Vaccine**, which can be explained: Tweets from both labels are expected to have many common features ("vaccine", "virus" etc.). The number of **Pro-Vaccine** tweets in the train set is significantly greater, which confuses the model to increase these feature weights on the **Pro-Vaccine** class. We can improve the model performance on **Anti-Vaccine** tweets, by "feeding" it with more of them.
- When we do not specify `min_df` & `max_df` parameters, the model is overfitting, since the number of features is way too large and prevents our model predictions from generalizing. In these cases, we can see that all the metric scores are better, which means that, hypothetically, if we feed the model with more samples, it will achieve better scores (when the two curves eventually approach each other). Of course, the difference between the two curves is huge, therefore the required amount of extra samples may be extreme.
- We see that preprocessing, as well as using stop words list does not help the model. In fact it causes a slight performance reduction, and also increases the "sense" of possible overfitting, since the curves become almost parallel at the end.
- Simply using `CountVectorizer` causes slight scores reduction, so `TfidfVectorizer` (which is essentially `CountVectorizer` & `TfidfTransformer`) seems to be the right choice.
- The final model that is selected in the notebook uses `TfidfVectorizer`, with custom `min_df`, `max_df` & `ngram_range` parameters and a Classifier with just `max_iter = 1000` and `multiclass = "multinomial"` parameters.

## Development

The notebook has been developed in WSL Ubuntu 20.04, using Visual Studio Code & Python 3.8.10. It has been tested successfully in Google Colab environment as well.